

A Practical Preimage Attack on the 19-Round SHA-256 Compression Function via a Global σ_0 -Difference Table

Kameldip Singh Basra

Independent Researcher

kameldipbasra@gmail.com

ORCID: 0009-0001-0977-4614

September 2026

Abstract

We present a practical preimage attack on the compression function of SHA-256 reduced to 19 of its 64 rounds, in the standard setting: the chaining value is the fixed initial value, the target is an arbitrary 256-bit digest, and the attacker chooses the 512-bit message block. The attack combines a backward chain that fixes eight state words from the target, a precomputed 16 GB table inverting $u \mapsto \sigma_0(u) - u$, and a two-stage fixed-point iteration over the three schedule-consistency constraints. Each attempt sweeps one 32-bit variable, costs $O(2^{32})$ table lookups, and runs in about 25 s on one NVIDIA H100. Six preimages of digests of random messages were found and independently verified: three in dedicated runs that succeeded after 3, 20 and 34 attempts, and three in a preregistered campaign of 240 attempts that gives the first direct estimate of the per-attempt success rate, $\hat{\mu} = 3/240 = 1.25 \times 10^{-2}$ (95% Poisson interval $[2.6 \times 10^{-3}, 3.7 \times 10^{-2}]$; all three events fell in the screened arm and the control arm gave 0/120). Attempts are far from interchangeable: in a paired campaign with matched exposure, a screening pass costing $1/4096$ of an attempt raised the yield of the intermediate 16-bit filter by $3.41 \times$ (95% CI $[2.13, 5.82]$, target-blocked bootstrap), helping on 31 of 40 targets and hurting on 9; its effect on the final success rate is not established.

SAT-based cryptanalysis has produced 17- and 18-round preimages of the zero digest, and a parameterized Cube-and-Conquer solver has recently been reported to invert the 19-round compression function. We make no priority claim at 19 rounds. The contribution here is a different route to the same round count, algebraic rather than search-based, with an explicit and low cost profile and complete artefacts. A companion note explains why the construction stops at 20 rounds. Full SHA-256 is not affected by any of these results.

Keywords: SHA-256, preimage attack, step-reduced, message schedule, table lookup, GPU cryptanalysis.

1 Introduction

SHA-256 [2] is the most widely deployed member of the SHA-2 family. Its compression function has 64 rounds, and the security margin of the full function is usually assessed through attacks on step-reduced variants.

Attack model. Let compress_R denote the SHA-256 compression function restricted to its first R rounds and let IV be the standard initial value. Given a 256-bit digest h , a *preimage* is a 512-bit block $(W_0, \dots, W_{15})$ with

$$h = \text{IV} + \text{compress}_R(\text{IV}, W_0, \dots, W_{15}).$$

This is the standard preimage setting for the compression function: the chaining value is fixed, nothing about any original message is known, and only the digest is given. It is not a pseudo-preimage or free-start setting, in which the attacker would also choose the chaining value. The 16 words of the block are unconstrained, so the result concerns the compression function and not the length-padded hash. For a single-block 55-byte message the padding fixes W_{14} , W_{15} and the low byte of W_{13} ; those three conditions can be met by choosing three of the attack’s context words in closed form (`padded_preimage.py` in the artefacts), but a 167-attempt run produced no padded preimage where about one was expected, so we claim only the compression-function result.

Prior work. Theoretical preimage attacks on step-reduced SHA-256 reach far beyond 19 rounds, but at costs within a few bits of the generic 2^{256} : Isobe and Shibutani [3] give a one-block preimage attack on 24 steps at 2^{240} ; Aoki et al. [4] reach 43 steps at $2^{254.9}$ by meet-in-the-middle; Khovratovich, Rechberger and Savelieva [5] reach 45 steps at $2^{255.5}$ with bicliques. None of these has been executed. Collision attacks are a different problem and are far more advanced, with practical collisions for 31 steps [8, 9] building on local-collision techniques [6, 7]; they do not yield preimages.

The practical preimage frontier belongs to SAT-based cryptanalysis [10]. Davydov, Pikhtovnikov, Kiryanova and Zaikin [11] find preimages of the 17- and 18-round compression function with Kissat under four CNF encodings, report 19 rounds as unsolved under every encoding, and reach 19 rounds only in a weakened form in which at most 31 of the 256 output bits are constrained; their published benchmark instances target the all-zero digest. Zaikin [12] then constructs intermediate inverse problems between consecutive round counts, parameterizes a CDCL solver on them, and reports that the resulting parallel Cube-and-Conquer solver inverted the 29-step MD5, 24-round SHA-1 and 19-round SHA-256 compression functions for the first time. The intermediate problems are weakened by construction; the abstract’s claim is for the compression function itself, and since the same group had already published the weakened 19-round result [11], “for the first time” can only refer to the full inversion. That article appeared in January 2026. We have not been able to examine its full text, so we cannot confirm that all 256 output bits were fixed, nor compare the attack model or the computational cost, and we make no priority claim at 19 rounds.

This work. We give an algebraic, table-based preimage attack on the 19-round compression function. Its cost profile is explicit and low: after a one-time 16 GB precomputation, each attempt is a single 2^{32} sweep on one GPU, about 25 seconds on an NVIDIA H100, and in the campaign of \$5.7 about one attempt in eighty succeeded. The ingredients are:

1. a backward chain that recovers the eight state words $a_{11}, \dots, a_{18}$ from the digest alone (§3);
2. a global table inverting $u \mapsto \sigma_0(u) - u$, built once and valid for every target and every attempt (§4);
3. the observation that in each of the three schedule constraints one unknown state word appears with opposite signs in two places, so that the constraint reduces to a single table lookup once the remaining unknowns are provisionally fixed (§5); the provisional assignment is closed by a fixed-point iteration and tested exactly by a 32-bit consistency check;
4. six independently verified preimages of digests of random messages, three from dedicated runs and three from a preregistered 240-attempt campaign that also measures the success rate directly (§6);

5. a measured, preregistered account of how much a cheap screening pass improves the intermediate yield, and of what it does not establish (§5.7).

A companion note [1] explains why the same construction does not reach 20 rounds. Appendix A records an exact cancellation identity in the message schedule found along the way, together with the reason the search experiments built on it are not claimed as results.

2 SHA-256 Preliminaries

Notation. All arithmetic is modulo 2^{32} unless stated otherwise. We use $\oplus$, $\&$, $\bar{x}$ for bitwise XOR, AND, NOT. Rotations and shifts of a 32-bit word x :

$$\text{ROTR}^n(x) = (x \gg n) \mid (x \ll (32 - n)), \quad \text{SHR}^n(x) = x \gg n.$$

Round functions.

$$\begin{aligned} \Sigma_0(a) &= \text{ROTR}^2(a) \oplus \text{ROTR}^{13}(a) \oplus \text{ROTR}^{22}(a), \\ \Sigma_1(e) &= \text{ROTR}^6(e) \oplus \text{ROTR}^{11}(e) \oplus \text{ROTR}^{25}(e), \\ \sigma_0(x) &= \text{ROTR}^7(x) \oplus \text{ROTR}^{18}(x) \oplus \text{SHR}^3(x), \\ \sigma_1(x) &= \text{ROTR}^{17}(x) \oplus \text{ROTR}^{19}(x) \oplus \text{SHR}^{10}(x), \\ \text{Ch}(e, f, g) &= (e \& f) \oplus (\bar{e} \& g), \\ \text{Maj}(a, b, c) &= (a \& b) \oplus (a \& c) \oplus (b \& c). \end{aligned}$$

State labels. We write the 8-word compression state after round r as $(a_r, a_{r-1}, a_{r-2}, a_{r-3}, e_r, e_{r-1}, e_{r-2}, e_{r-3})$, with IV boundary values $a_{-4}, \dots, a_{-1}$ and $e_{-4}, \dots, e_{-1}$:

$$a_{-1}=\text{IV}[0], \quad a_{-2}=\text{IV}[1], \quad a_{-3}=\text{IV}[2], \quad a_{-4}=\text{IV}[3], \quad e_{-1}=\text{IV}[4], \quad e_{-2}=\text{IV}[5], \quad e_{-3}=\text{IV}[6], \quad e_{-4}=\text{IV}[7].$$

Round equations.

$$T_1^{(r)} = \Sigma_1(e_{r-1}) + \text{Ch}(e_{r-1}, e_{r-2}, e_{r-3}) + e_{r-4} + K_r + W_r, \quad (1)$$

$$T_2^{(r)} = \Sigma_0(a_{r-1}) + \text{Maj}(a_{r-1}, a_{r-2}, a_{r-3}), \quad (2)$$

$$a_r = T_1^{(r)} + T_2^{(r)}, \quad (3)$$

$$e_r = a_{r-4} + T_1^{(r)}. \quad (4)$$

Note that e_{r-4} in (1) is the *h-register* word (boundary value $e_{-4} = \text{IV}[7]$), while a_{r-4} in (4) is the *d-register* word (boundary value $a_{-4} = \text{IV}[3]$); these are distinct IV words and distinct chain values for all $r \geq 4$.

From (3)–(4): $T_1^{(r)} = a_r - T_2^{(r)}$, and rearranging (1):

$$W_r = T_1^{(r)} - \Sigma_1(e_{r-1}) - \text{Ch}(e_{r-1}, e_{r-2}, e_{r-3}) - e_{r-4} - K_r. \quad (5)$$

Message schedule. For $r \geq 16$ the message schedule expands:

$$W_r = \sigma_1(W_{r-2}) + W_{r-7} + \sigma_0(W_{r-15}) + W_{r-16}. \quad (6)$$

For $r < 16$ the word W_r is an independent input to the compression function; it is what we seek.

Hash output. After R rounds the hash is $h[i] = \text{IV}[i] + s_i$ for $i = 0, \dots, 7$, where the final state vector is $(s_0, \dots, s_7) = (a_{R-1}, a_{R-2}, a_{R-3}, a_{R-4}, e_{R-1}, e_{R-2}, e_{R-3}, e_{R-4})$.

3 Structural Analysis

3.1 Backward Chain

From the R -round hash we recover the final state directly:

$$a_{R-1}, a_{R-2}, a_{R-3}, a_{R-4}, e_{R-1}, e_{R-2}, e_{R-3}, e_{R-4}.$$

We then extend backward using $T_2^{(r)} = \Sigma_0(a_{r-1}) + \text{Maj}(a_{r-1}, a_{r-2}, a_{r-3})$, $T_1^{(r)} = a_r - T_2^{(r)}$, $a_{r-4} = e_r - T_1^{(r)}$, iterating for $r = R-1, R-2, R-3, R-4$ to recover $a_{R-5}, \dots, a_{R-8}$.

For $R = 19$ this yields $\{a_{11}, a_{12}, \dots, a_{18}\}$: the state words a_{11} through a_{18} are determined by the target hash alone, with no knowledge of any preimage.

3.2 Free Parameters and Schedule Constraints

At $R = 19$ the unconstrained parameters are $a_0, \dots, a_{10}$ (11 words). Of these, $a_4, \dots, a_{10}$ are chosen uniformly at random per attempt (*context*), forming 7 attacker-controlled free words. The remaining $a_0, \dots, a_3$ are the targets of the attack phases.

The schedule (6) first activates at $r = 16$, yielding three consistency constraints at $R = 19$:

$$W_{16} = \sigma_1(W_{14}) + W_9 + \sigma_0(W_1) + W_0, \quad (\text{C0})$$

$$W_{17} = \sigma_1(W_{15}) + W_{10} + \sigma_0(W_2) + W_1, \quad (\text{C1})$$

$$W_{18} = \sigma_1(W_{16}) + W_{11} + \sigma_0(W_3) + W_2. \quad (\text{C2})$$

Here W_{16}, W_{17}, W_{18} are recovered from the target alone via the backward chain and (5), using the known $a_{11}, \dots, a_{18}$ together with the context values $a_4, \dots, a_{10}$ (which supply $e_{11}, \dots, e_{14}$, needed to evaluate (5) for $r = 16, 17, 18$). Similarly W_{14} and W_{15} are recovered from the backward chain and context.

4 The σ_0 -Difference Table

Definition 1. The σ_0 -difference table $\mathcal{T}$ is an array of 2^{32} 32-bit entries:

$$\mathcal{T}[v] = u$$

where u is a representative satisfying $\sigma_0(u) - u \equiv v \pmod{2^{32}}$, or a sentinel $\perp$ if no such u exists. When multiple u share the same v , one representative is stored.

Construction. $\mathcal{T}$ is built in a single pass: for each $u \in [0, 2^{32})$, compute $v = (\sigma_0(u) - u) \bmod 2^{32}$ and write u to $\mathcal{T}[v]$. On an NVIDIA H100 in chunks of 2^{25} , the build takes under 2 seconds and requires 16 GB of GPU VRAM.

Coverage. The function $f(u) = \sigma_0(u) - u \bmod 2^{32}$ is not injective. Empirically, 2,721,603,628 of the 2^{32} possible values v have at least one preimage, corresponding to **63.37%** coverage.

Usage. Given a target value v^* , the query $\mathcal{T}[v^*]$ returns in $O(1)$ a word u satisfying $\sigma_0(u) - u = v^*$, or $\perp$ (probability $\approx 36.6\%$).

5 The 19-Round Preimage Attack

5.1 Overview

The attack operates in three phases per context:

C0. Sweep $a_0 \in [0, 2^{32})$. A table lookup recovers W_1 (and hence a_1) for 63.37% of a_0 values.

C1/C2. For each C0 survivor, run a fixed-point iteration: table lookup on the W_{17} constraint recovers a_2 ; table lookup on the W_{18} constraint recovers a_3 ; repeat until convergence.

W9 filter. Test the consistency condition $\varepsilon(a_2, a_3) = 0$ of Proposition 1 exactly; its pass rate is measured in §5.6. Candidates that pass are verified by forward evaluation.

The key that makes all three phases work is a common algebraic pattern: in each constraint, the target unknown a_i appears with *opposite signs* in two different terms, causing exact cancellation in the σ_0 -difference lookup target.

5.2 Context Precomputation

Fix a target hash h . Run the backward chain to obtain $a_{11}, \dots, a_{18}$. Choose uniformly random context $a_4, \dots, a_{10}$. Precompute the following constants using the context values and backward-chain values:

$$T_2^{\text{IV}} = \Sigma_0(a_{-1}) + \text{Maj}(a_{-1}, a_{-2}, a_{-3}), \quad (7)$$

$$C_0 = -T_2^{\text{IV}} - e_{-4} - \Sigma_1(e_{-1}) - \text{Ch}(e_{-1}, e_{-2}, e_{-3}) - K_0, \quad (8)$$

$$C_{e_0} = a_{-4} - T_2^{\text{IV}}. \quad (9)$$

From these: $W_0 = a_0 + C_0$ and $e_0 = a_0 + C_{e_0}$, both linear in a_0 .

For the context rounds $r = 7, \dots, 11$, compute $T_2^{(r)}, T_1^{(r)}, e_r$ from (1)–(4) using the context $a_4, \dots, a_{10}$ together with a_{11} from the backward chain. Define:

$$W_9^{\text{base}} = T_1^{(9)} - \Sigma_1(e_8) - K_9, \quad (10)$$

$$W_{11}^{\text{base}} = T_1^{(11)} - T_1^{(7)} - \Sigma_1(e_{10}) - \text{Ch}(e_{10}, e_9, e_8) - K_{11}, \quad (11)$$

$$\text{CONST}_{10} = T_1^{(10)} - \Sigma_1(e_9) - K_{10}. \quad (12)$$

Also recover $W_{14}, W_{15}, W_{16}^{\text{bc}}, W_{17}^{\text{bc}}, W_{18}^{\text{bc}}$ from (5) at rounds 14, 15, 16, 17, 18 respectively (all from the backward chain and the context).

Finally, initialise $\hat{W}_9$ with $a_2 = a_3 = 0$:

$$\hat{W}_9 = W_9^{\text{base}} - T_1^{(5)}|_{a_3=0} - \text{Ch}\left(e_8, T_1^{(7)} + 0, T_1^{(6)}|_{a_2=a_3=0}\right), \quad (13)$$

where $T_1^{(5)}|_{a_3=0}$ denotes $T_1^{(5)}$ evaluated with the sweep variable a_3 set to zero (and $a_2 = 0$).

5.3 Phase C0: The Provisional- W_9 Form

The constraint C0 involves W_9 , which depends on the unknowns a_1, a_2, a_3 . The attack does not solve C0 exactly. It solves C0 under a provisional assignment for the part of W_9 that depends on a_2, a_3 , and tests that assignment afterwards. The following proposition makes the assignment precise.

Proposition 1 (Provisional- W_9 form). *Let W_9 be the message word determined by the state through equation (5) at $r = 9$, and let $\hat{W}_9$ be the quantity (13), computed from the context with $a_2 = a_3 = 0$. Then*

$$\hat{W}_9 - W_9 = a_1 + \varepsilon(a_2, a_3) \pmod{2^{32}},$$

where

$$\varepsilon(a_2, a_3) = [T_1^{(5)} - T_1^{(5)}|_{a_2=a_3=0}] + [\text{Ch}(e_8, e_7, e_6) - \text{Ch}(e_8, T_1^{(7)}, T_1^{(6)}|_{a_2=a_3=0})]$$

depends on a_2, a_3 only through $\text{Maj}(a_4, a_3, a_2)$, $\text{Maj}(a_5, a_4, a_3)$ and the Ch arguments $e_6 = a_2 + T_1^{(6)}$ and $e_7 = a_3 + T_1^{(7)}$, and vanishes at $a_2 = a_3 = 0$.

Proof. Apply (5) at $r = 9$:

$$W_9 = \underbrace{T_1^{(9)} - \Sigma_1(e_8) - K_9}_{W_9^{\text{base}}} - e_5 - \text{Ch}(e_8, e_7, e_6).$$

$T_1^{(9)}$ and e_8 depend only on the context $a_4, \dots, a_9$, so W_9^{base} is a context constant. The identity $e_r = a_{r-4} + T_1^{(r)}$ at $r = 5, 6, 7$ gives

$$e_5 = a_1 + T_1^{(5)}, \tag{14}$$

$$e_6 = a_2 + T_1^{(6)}, \tag{15}$$

$$e_7 = a_3 + T_1^{(7)}, \tag{16}$$

where $T_1^{(7)} = a_7 - T_2^{(7)}$ is a context constant ($T_2^{(7)} = \Sigma_0(a_6) + \text{Maj}(a_6, a_5, a_4)$), $T_1^{(6)}$ depends on a_3 through $\text{Maj}(a_5, a_4, a_3)$, and $T_1^{(5)}$ depends on a_2, a_3 through $\text{Maj}(a_4, a_3, a_2)$. Substituting (14),

$$W_9 = W_9^{\text{base}} - a_1 - T_1^{(5)} - \text{Ch}(e_8, e_7, e_6),$$

while by definition (13)

$$\hat{W}_9 = W_9^{\text{base}} - T_1^{(5)}|_{a_2=a_3=0} - \text{Ch}(e_8, T_1^{(7)}, T_1^{(6)}|_{a_2=a_3=0}).$$

Subtracting gives the stated expression. The two brackets are the a_2, a_3 -dependent parts of $T_1^{(5)}$ and of the Ch term, and both are zero at $a_2 = a_3 = 0$. $\square$

The residual ε is not zero in general: on a random state it is a nonzero 32-bit word. The attack proceeds *as if* $\varepsilon = 0$. We call this the *provisional assignment*. Under it, C0 reduces to a single table lookup (Proposition 2); C1 and C2 are then solved for a_2, a_3 by iteration (§5.4); and finally $\varepsilon(a_2, a_3) = 0$ is tested exactly (§5.5). A candidate that passes the test satisfies all three constraints with the true W_9 , and forward evaluation confirms it. The attack therefore finds those preimages whose (a_2, a_3) happen to satisfy $\varepsilon = 0$. The reason this is a productive subset, rather than a 2^{-32} fraction of the solutions, is that the iteration is seeded on partial consistency with $\hat{W}_9$ and converges preferentially to such points; how often is an empirical question, answered in §5.6 by measurement. For each of the six verified preimages $\hat{W}_9 - W_9 = a_1$ holds exactly, as it must for any solution the attack returns; for P1, $0x72c05667 - 0x6d1e5a20 = 0x05a1fc47 = a_1$.

C0 lookup. For each $a_0 \in [0, 2^{32})$, compute:

$$e_0 = a_0 + C_{e_0}, \tag{17}$$

$$W_0 = a_0 + C_0, \tag{18}$$

$$g(a_0) = -\Sigma_0(a_0) - \text{Maj}(a_0, a_{-1}, a_{-2}) - e_{-3} - \Sigma_1(e_0) - \text{Ch}(e_0, e_{-1}, e_{-2}) - K_1, \tag{19}$$

$$F_0 = W_{16}^{\text{bc}} - \sigma_1(W_{14}) - g(a_0) - \hat{W}_9 - W_0. \tag{20}$$

Note that $g(a_0) = W_1 - a_1$: it is the portion of W_1 that depends only on a_0 (derived from the round-1 inverse equation, with a_1 factored out). Then:

Proposition 2 (C0 cancellation). *Under the provisional assignment $\varepsilon(a_2, a_3) = 0$, $F_0 = \sigma_0(W_1) - W_1 \pmod{2^{32}}$.*

Proof. Under the provisional assignment, Proposition 1 gives $W_9 = \hat{W}_9 - a_1$, so from the C0 schedule constraint

$$\sigma_0(W_1) = W_{16}^{\text{bc}} - \sigma_1(W_{14}) - W_9 - W_0 = W_{16}^{\text{bc}} - \sigma_1(W_{14}) - \hat{W}_9 + a_1 - W_0.$$

Since $W_1 = a_1 + g(a_0)$:

$$\sigma_0(W_1) - W_1 = W_{16}^{\text{bc}} - \sigma_1(W_{14}) - \hat{W}_9 + a_1 - W_0 - a_1 - g(a_0) = F_0. \quad \square$$

The a_1 terms cancel exactly. A single table lookup $\mathcal{T}[F_0]$ returns W_1 , and then $a_1 = W_1 - g(a_0)$.

5.4 Phases C1/C2: Fixed-Point Iteration

For each C0 survivor (a_0, a_1, W_1) , derive:

$$e_1 = e_{-3} + a_1 - \Sigma_0(a_0) - \text{Maj}(a_0, a_{-1}, a_{-2}), \quad (21)$$

$$T_2^{(2)} = \Sigma_0(a_1) + \text{Maj}(a_1, a_0, a_{-1}), \quad (22)$$

$$F_{12} = -T_2^{(2)} - e_{-2} - \Sigma_1(e_1) - \text{Ch}(e_1, e_0, e_{-1}) - K_2. \quad (23)$$

Note that $F_{12} = W_2 - a_2$: the portion of W_2 depending only on (a_0, a_1) .

Proposition 3 (C1 and C2 cancellations). *Define the a_3 -dependent quantity:*

$$D(a_3) = (a_6 - \Sigma_0(a_5) - \text{Maj}(a_5, a_4, a_3)) + \text{Ch}(e_9, e_8, a_3 + T_1^{(7)}).$$

Let $W_{16}^{\text{sched}} = \sigma_1(W_{14}) + \hat{W}_9 + \sigma_0(W_1) + W_0$ and $W_{16}^* = W_{16}^{\text{sched}} - a_1$. Then:

$$F_{C1}(a_3) = W_{17}^{\text{bc}} - \sigma_1(W_{15}) - W_1 - \text{CONST}_{10} - F_{12} + D(a_3) = \sigma_0(W_2) - W_2, \quad (\text{C1-target})$$

$$F_{C2}(a_2, a_3) = W_{18}^{\text{bc}} - \sigma_1(W_{16}^*) - W_{11}^{\text{base}} - W_2 - F_{23}(a_0, a_1, a_2) = \sigma_0(W_3) - W_3, \quad (\text{C2-target})$$

where $F_{23}(a_0, a_1, a_2)$ is the portion of W_3 depending on (a_0, a_1, a_2) but not a_3 . The C1 identity is unconditional; the C2 identity holds under the provisional assignment of Proposition 1, which enters through W_{16}^* .

Proof. C1. From (5) at $r = 10$, using $\text{CONST}_{10} = T_1^{(10)} - \Sigma_1(e_9) - K_{10}$:

$$W_{10} = \text{CONST}_{10} - e_6 - \text{Ch}(e_9, e_8, e_7).$$

By the shift-register identity $e_r = a_{r-4} + T_1^{(r)}$: $e_6 = a_2 + T_1^{(6)}(a_3)$ and $e_7 = a_3 + T_1^{(7)}$, so $\text{Ch}(e_9, e_8, e_7) = \text{Ch}(e_9, e_8, a_3 + T_1^{(7)})$. Grouping the a_3 -dependent terms into $D(a_3)$:

$$W_{10} = \text{CONST}_{10} - a_2 - D(a_3).$$

Substituting into the C1 schedule constraint $\sigma_0(W_2) = W_{17}^{\text{bc}} - \sigma_1(W_{15}) - W_{10} - W_1$:

$$\begin{aligned} \sigma_0(W_2) &= W_{17}^{\text{bc}} - \sigma_1(W_{15}) - (\text{CONST}_{10} - a_2 - D(a_3)) - W_1 \\ &= W_{17}^{\text{bc}} - \sigma_1(W_{15}) - \text{CONST}_{10} + a_2 + D(a_3) - W_1. \end{aligned}$$

Using $W_2 = a_2 + F_{12}$, the a_2 terms cancel:

$$\begin{aligned} \sigma_0(W_2) - W_2 &= (W_{17}^{\text{bc}} - \sigma_1(W_{15}) - \text{CONST}_{10} + a_2 + D(a_3) - W_1) - (a_2 + F_{12}) \\ &= W_{17}^{\text{bc}} - \sigma_1(W_{15}) - W_1 - \text{CONST}_{10} - F_{12} + D(a_3) = F_{C1}(a_3). \quad \square_{C1} \end{aligned}$$

C2. From (5) at $r = 11$: $W_{11} = T_1^{(11)} - \Sigma_1(e_{10}) - \text{Ch}(e_{10}, e_9, e_8) - e_7 - K_{11}$. Since $e_7 = a_3 + T_1^{(7)}$:

$$W_{11} = \underbrace{[T_1^{(11)} - T_1^{(7)} - \Sigma_1(e_{10}) - \text{Ch}(e_{10}, e_9, e_8) - K_{11}]}_{W_{11}^{\text{base}}} - a_3.$$

So $W_{11} = W_{11}^{\text{base}} - a_3$ with W_{11}^{base} determined by context alone.

Under the provisional assignment (Proposition 1), $W_{16} = W_{16}^{\text{sched}} - a_1 = W_{16}^*$, so $\sigma_1(W_{16}) = \sigma_1(W_{16}^*)$. Substituting into C2: $\sigma_0(W_3) = W_{18}^{\text{bc}} - \sigma_1(W_{16}^*) - W_{11} - W_2 = W_{18}^{\text{bc}} - \sigma_1(W_{16}^*) - (W_{11}^{\text{base}} - a_3) - W_2$. Using $W_3 = a_3 + F_{23}$, the a_3 terms cancel:

$$\begin{aligned} \sigma_0(W_3) - W_3 &= W_{18}^{\text{bc}} - \sigma_1(W_{16}^*) - W_{11}^{\text{base}} + a_3 - W_2 - a_3 - F_{23} \\ &= W_{18}^{\text{bc}} - \sigma_1(W_{16}^*) - W_{11}^{\text{base}} - W_2 - F_{23} = F_{C2}(a_2, a_3). \end{aligned}$$

□

Iteration. The joint C1/C2 system defines a map

$$(a_2, a_3) \mapsto (\mathcal{T}[F_{C1}(a_3)] - F_{12}, \mathcal{T}[F_{C2}(a_2, a_3)] - F_{23}).$$

We iterate from a seed $(a_2^{(0)}, a_3^{(0)})$ for up to 8 steps. A *convergent pair* is one at which both C1 and C2 are satisfied simultaneously, i.e., the map is a fixed point. Empirical convergence rate: $1.3\text{--}1.9 \times 10^{-5}$ per C0 survivor (E1, §5.6).

K-seed optimisation. Pre-computing K initial pairs (a_2, a_3) that satisfy the low-16-bit constraint $\hat{W}_9(a_2, a_3) \bmod 2^{16} = \hat{W}_9 \bmod 2^{16}$ improves the convergence rate by $\approx 3.3\times$.

5.5 W9 Bit-Filter

After each convergent pair (a_2, a_3) is found, recompute $\hat{W}_9 - \varepsilon(a_2, a_3)$, the a_1 -free part of W_9 at the converged values:

$$W_9^{\text{conv}} = W_9^{\text{base}} - (a_5 - \Sigma_0(a_4) - \text{Maj}(a_4, a_3, a_2)) - \text{Ch}(e_8, a_3 + T_1^{(7)}, a_2 + T_1^{(6)}(a_3)). \quad (24)$$

Two sequential filters are applied:

- **Lo filter:** $(W_9^{\text{conv}} \bmod 2^{16}) = (\hat{W}_9 \bmod 2^{16})$, i.e. the low half of ε vanishes. Passes at $52\text{--}85\times$ (E1) or $37\text{--}44\times$ (E2) the uniform 2^{-16} rate (experiments defined in §5.6).
- **Hi filter:** $(W_9^{\text{conv}} \gg 16) = (\hat{W}_9 \gg 16)$, i.e. the high half vanishes too, so that both together test $\varepsilon = 0$ exactly. Passes at $12.6\text{--}13.3\times$ (E1) or $7\text{--}11\times$ (E2) the uniform rate. This filter is *not* uniform; a coarse-binned test suggests otherwise but is underpowered against a spike at a single residual (§6.3).

The combined rate is not obtained by multiplying the two marginals: the iteration is seeded on lo-consistency, so independence of the two halves cannot be assumed and has not been established (§5.6). Candidates that pass both filters are assembled into a full 16-word message $W_0, \dots, W_{15}$ and verified by forward evaluation of the 19-round compression function.

Table 1: Empirical per-context statistics. E1: two deduplicated CPU pilots of 0.20 and 1.37 context-equivalents. Lo-pass counts per context: the E1 pilots, the 165-context H100 run of §6.3, and the campaign control arm of §5.7. μ rows: the 240-context campaign. Wall times: H100.

Quantity	Value
C0 survivors (table hits)	2,721,603,628 (63.37% of 2^{32})
E1 — <i>deduplicated (unique fixed points); two CPU pilots</i>	
C1/C2 fixed points	$3.6\text{--}5.1 \times 10^4$ per context ($1.3\text{--}1.9 \times 10^{-5}$ per C0 survivor)
W_9 lo-pass events per context	28 (E1), 35.0 (165-context run), 33.7 (campaign control)
W_9 lo-pass rate per fixed point	$7.9 \times 10^{-4}\text{--}1.3 \times 10^{-3}$
Structural <i>lo</i> enhancement	$52\text{--}85\times$ above 2^{-16}
W_9 hi-pass rate per fixed point	$\approx 2.0 \times 10^{-4}$
Structural <i>hi</i> enhancement	$12.6\text{--}13.3\times$ above 2^{-16}
<i>Legacy trajectory counts (each fixed point counted once per remaining iteration)</i>	
Repeat-count inflation factor	$5.5\text{--}5.9\times$
μ (original model: measured lo count $\times$ uniform conditional hi)	
	$\approx 5 \times 10^{-4}$
μ (directly measured, 240 ctx)	1.25×10^{-2} CI [$2.6\cdot 10^{-3}, 3.7\cdot 10^{-2}$]
screened arm (120 ctx, $N_{LH} = 3$)	2.50×10^{-2}
control arm (120 ctx, $N_{LH} = 0$)	$< 3.1 \times 10^{-2}$ (exact Poisson, 95%)
μ (independence-derived projection)	$6 \times 10^{-3}\text{--}1.3 \times 10^{-2}$
Observed contexts (P3, P2, P1)	3, 20, 34 (mean 19)
Wall time per context (H100), $K_{\text{seeds}}=0$	≈ 25 s
same, $K_{\text{seeds}}=4$ (campaign)	≈ 89 s

5.6 Complexity Analysis

W9 structural enhancement. Neither half of the W_9 filter is uniform. Both enhancements arise because the C1/C2 fixed-point iteration, started from W_9 -lo-consistent seed pairs, preferentially converges to fixed points that also satisfy the W_9 constraint—an effect that acts on the full 32-bit residual, not only on the low half. Experiment **E2**, an independent CPU reimplementing measuring over 2.2×10^9 swept a_0 values per configuration, gives:

- the *lo* filter passes at $37\text{--}44\times$ the uniform 2^{-16} rate, confirming the enhancement mechanism;
- the *hi* filter passes at $7\text{--}11\times$ the uniform rate—it is *not* uniform, contrary to what a coarse-binned test suggests (§6.3).

E1 and E2 are separate campaigns and their rates are not directly comparable: E1 deduplicates fixed points and reports per-fixed-point rates from two CPU pilots, whereas E2 reports per-swept- a_0 rates from a distinct implementation on different contexts. Both find the same qualitative effect — a large *lo* enhancement and a smaller but nonzero *hi* enhancement — and the residual factor of roughly 1.5 between them is not currently explained; it is within the spread expected from context-to-context variation (§5.7), which is itself large. **E1 is the headline dataset**, since it is deduplicated and feeds the μ projection. Neither E1 nor E2 measures the success rate; the preregistered campaign below does.

These rates are not artefacts of the lookup table’s representative policy. A third experiment (**E3**) repeated the measurement on byte-identical contexts with a table built to keep the *smallest* rather than the *largest* preimage of each target: lo enhancement $30.0\times$ (smallest) against $28.5\times$ (largest), and hi $11.1\times$ against $6.9\times$, agreeing within Poisson error on the 8 and 13 hi events observed. E3 used its own contexts and sample, which is why its lo figure sits below E1’s range. Direct tests of lo/hi independence at reduced filter widths give $D = P(\text{both})/[P(\text{lo})P(\text{hi})] =$

1.00 ± 0.03 at 4 bits and 1.06 ± 0.10 at 6 bits, consistent with independence at the widths where joint events can be counted.

Multiplying the measured marginals by the deduplicated fixed-point count gives a *projection* of $\mu \approx 6 \times 10^{-3}$ – 1.3×10^{-2} per context (one figure per pilot). Once fixed points are deduplicated the lo-pass count is 28–35 per context across the pilots, the 165-context run and the campaign control arm (Table 1), so the lo analysis stands; the correction is almost entirely the hi filter, worth about $13\times$ rather than the two orders of magnitude a naive comparison of non-deduplicated counts suggests. We label this a projection rather than a measurement for two reasons. First, it assumes the two halves are independent; since the iteration is seeded on lo-consistency, that is exactly the assumption one should not make for free. Half-width diagnostics detect no material dependence at 4 and 6 bits, but a sparse spike at the exact 16-bit value would remain invisible to them—the same limitation that makes the binned χ^2 test uninformative (§6.3). Second, the direct estimator $\hat{\mu} = N_{LH}/C$ now has events to work with, though few. The preregistered campaign of §5.7 attacked 240 contexts at full 2^{32} exposure and recorded $N_{LH} = 3$ complete 32-bit matches, all of which verified as preimages, giving

$$\hat{\mu} = \frac{3}{240} = 1.25 \times 10^{-2} \text{ per context,} \quad 95\% \text{ CI } [2.58 \times 10^{-3}, 3.65 \times 10^{-2}].$$

Split by arm, the screened half gave $3/120 = 2.50 \times 10^{-2}$ and the control half $0/120$, i.e. $\mu < 3.1 \times 10^{-2}$ (upper limit of the exact two-sided 95% Poisson interval for zero events in 120 contexts). Two cautions apply. The pooled figure mixes the two arms and the screened half is enriched by construction, so it is not an estimate for uniformly sampled contexts; the control arm is the closer analogue and it produced no events. And three events give an interval spanning a factor of 14, so this bounds the rate rather than pinning it. What it does settle is that the rate is not pathologically below the projection: the measured 1.25×10^{-2} falls inside the 6×10^{-3} – 1.3×10^{-2} band projected before the campaign was run, which is a genuine out-of-sample check on the independence-derived model rather than a fit to it. The projections themselves remain unstable. After adopting global fixed-point deduplication, the independence-derived projection was 1.31×10^{-2} in a 0.20 context-equivalent pilot and 5.80×10^{-3} in a 1.37 context-equivalent pilot. That variation shows the marginal-event sample is too small for a stable projection; it does not by itself establish a directional bias. An earlier figure of 4.88×10^{-2} used repeat-counted fixed points and is not comparable with either.

Two earlier sources of error are worth recording, since both inflate μ if left uncorrected. A converged lane remains at its fixed point for every subsequent iteration, so counting convergences per iteration overcounts each fixed point—measured at $5.89\times$ for eight iterations. And convergences tallied per seed chain rather than per context understate the per-context count by the number of chains. All figures above are globally deduplicated on the fixed-point identity (a_0, a_1, a_2, a_3) within a context.

What the three dedicated finds do and do not establish. The three dedicated runs found preimages after 3, 20 and 34 contexts: three successes over 57 attempted contexts, $\hat{p} = 0.053$ with an exact 95% interval of $[0.011, 0.146]$. The 6×10^{-3} projection lies *below* that interval. Under $p = 0.006$, observing three or more successes in 57 contexts has probability 0.005, so either the historical runs were unusually favourable or the projection is low. Two very small samples are in tension here, and neither settles the rate. It is not compatible with the uniform-filter model: at $p \approx 5 \times 10^{-4}$ the same observation has probability 4×10^{-6} , so the finds do falsify the uniform model even though they cannot pin the true rate. We therefore present them as demonstrations of feasibility, not as a runtime estimator. Three observations cannot support an expected-value claim, and the per-context success probability $\Pr(N_{LH} \geq 1)$ —which governs time to first preimage—is not in general μ itself if successful fixed points cluster in favourable contexts. The preregistered multi-target campaign of §5.7 reports both: over 240

fully attacked contexts, $\hat{\mu} = 3/240 = 1.25 \times 10^{-2}$ and $\widehat{\Pr}(N_{LH} \geq 1) = 3/240$ also, no context having produced more than one match. With three events these coincide and neither is tightly determined. The combined search time for the three was 1,464 s on the H100 (871.6, 511.1 and 81.7 s), i.e. ≈ 25.7 s per context with the single (0,0) start. We note for comparison that the screening campaign (§5.7), which uses $K_{\text{seeds}}=4$ and therefore runs four seed chains per context, measured 89 s per context over 240 contexts. The ratio 3.5 is consistent with the $4\times$ increase in inner-loop work and not with any difference in the sweep itself; per-context time was flat in the fixed-point count (90.2 s in the screened arm against 88.3 s in the control arm despite $8.1\times$ the fixed points), confirming the cost is the 2^{32} sweep rather than the downstream bookkeeping.

5.7 Context Screening

The clustering just hypothesised is real and is large enough to exploit. The attack samples a context and then sweeps a_0 over 2^{32} , giving every context the same budget; that is optimal only if contexts are interchangeable. Over 280 contexts, each scored on a 2^{20} sample of a_0 , the deduplicated C1/C2 fixed-point count had mean 63.4 and variance 97,082, a variance-to-mean ratio of $\approx 1,530$ against the 1.0 that a Poisson model would give, with median 12 and maximum 4,924. The top decile of contexts held 68.6% of the total yield.

Overdispersion alone is not usable, since it could be per-sample noise. The operative question is whether a cheap score measured on one a_0 sample predicts yield on an independent one. Scoring each of 392 contexts twice on independent 2^{20} samples gave a Spearman rank correlation of 0.910 between the two scores, so productivity is a stable property of the context and is recoverable from a short pass. Since screening costs 2^{20} against a full sweep of 2^{32} , it is $1/4096$ of a context’s cost and is nearly free: selecting the top 1% on the screening sample gave a $12.99\times$ out-of-sample *fixed-point* yield lift for a 2.44% overhead. We stress that this is a statement about fixed points only; the attack-relevant lift is measured below and is substantially smaller. The deep tail is also unstable—the top-1% row rests on 4 contexts of 392 and moved from $12.99\times$ to $31.93\times$ between runs—so the top-decile figure ($4.98\times$ and $6.79\times$ in two runs) is the defensible one at this level.

Does the lift survive conversion? The score counts C1/C2 fixed points, so the lift above is in the first place a fixed-point lift. Whether it survives to the events that actually determine attack cost is a separate question, and we settled it with a preregistered paired campaign (`campaign_screening.py`, seed 20260810; the run manifest in `data_campaign_screening.json` records the command, seeds and software versions) rather than by assuming proportionality.

Design, fixed before the run. 40 independent targets; 10 contexts generated per target; every context scored on an identical cheap 2^{20} a_0 sample; the top 30% by score forming the SCREENED arm and an equal number drawn at random from the remainder forming the CONTROL arm; both arms then given *exactly* the same full 2^{32} exposure per context, with no early stop on a hit, so that a lucky context is never under-exposed. 240 contexts were fully attacked, 120 per arm, in 5.95 h of attack time on one H100 (screening excluded).

Counter (per context)	screened	control	lift
C1/C2 fixed points	159,545	19,624	$8.13\times$
W_9 lo-passes	114.86	33.73	$3.41\times$
W_9 hi-passes (given lo)	0.025	0	—
full 32-bit W_9	0.025	0	—
verified preimages	0.025	0	—

The conversion is sublinear. An $8.13\times$ lift in fixed points buys only a $3.41\times$ lift in lo-passes: screened contexts yield more fixed points, but each individual fixed point is *less* likely to pass

the lo filter. An earlier pilot on 10 lo events suggested the conversion was proportional; that pilot was underpowered and this campaign superseded it.

The effect is real but highly heterogeneous. A target-blocked bootstrap (20,000 resamples, blocking on target because contexts within a target are not independent) gives a 95% interval of $[2.13\times, 5.82\times]$ on the pooled lo lift. Per target the median lift is $3.78\times$ and the geometric mean $3.80\times$, but the interquartile range (Weibull quantiles) spans $1.31\times$ to $12.35\times$: screening helped on 31 of 40 targets and *hurt* on the other 9, where it selected contexts richer in fixed points but poorer in lo-passes. A distribution-free sign test on the per-target lifts—which is insensitive to the heavy tail that makes the pooled ratio unstable—gives $p = 3.4 \times 10^{-4}$ one-sided (6.8×10^{-4} two-sided).

What remains unmeasured. The screened arm produced 3 full 32-bit W_9 matches and 3 verified preimages against 0 in control. Under the null of equal rates each event falls in the screened arm with probability $1/2$, so all three doing so has $p = 1/8 = 0.125$ (one-sided): suggestive and *not* significant. We therefore report no lift for the full-match or preimage counters. Establishing one would require of order 10^2 GPU-hours, which we did not spend.

Summary of the claim. Screening costs $1/4096$ of a context’s budget and delivers a measured $3.41\times$ (95% CI $[2.13, 5.82]$) improvement in lo-pass yield per attacked context. It is **not** the $12.99\times$ figure quoted for fixed points, which we retain above only as a statement about fixed points; and it is not yet established as a speedup in time-to-preimage.

6 Experimental Results

6.1 Hardware and Software

All GPU experiments were run on NVIDIA H100 SXM5 80 GB HBM3 via RunPod cloud compute; the E1 and E2 pilots ran on CPU. The attack code uses PyTorch 2.x on CUDA, with the C0 sweep vectorised in batches of 2^{24} elements. The 16 GB σ_0 -difference table is held in GPU memory throughout the sweep; table lookups are implemented as indexed tensor reads.

6.2 Verified Preimages

Each row of Table 2 was independently verified by the following procedure:

1. Parse the 16 words $W_0, \dots, W_{15}$ as a 512-bit block.
2. Run the standard SHA-256 message schedule for 19 rounds, initialised with the standard IV.
3. Add the final state to the IV word-by-word.
4. Confirm the resulting 32-byte value equals the target hash.

All six preimages pass this check. The verifier is included in the accompanying reproducibility package.

6.3 Chi-Squared Uniformity Test

To test the $1/2^{16}$ model for the W9-hi filter, we collected the values $(W_9^{\text{conv}} \gg 16) \bmod 2^{16}$ for all lo-pass events across 165 GPU contexts (a total of 5,775 lo-pass events). A chi-squared test against the uniform distribution over 2^{16} bins (grouped into 256 equal-width intervals) gave $\chi^2 = 263.4$ with 255 degrees of freedom ($p = 0.35$).

This test is *not* sensitive to the effect that matters, and we record the point because it is easy to be misled by it. The enhancement is concentrated at and near the exact 16-bit residual, whereas grouping into 256 equal-width intervals spreads it across a bin of width 256. We

Table 2: Verified 19-round preimages. P1–P3 were found in dedicated runs ($K_{\text{seeds}}=0$); P4–P6 in the screened arm of the campaign of §5.7 ($K_{\text{seeds}}=4$; target and context number given). Only the target digest was supplied to the solver. All values are hexadecimal. [†]Found inside the third context.

	Hash / words				Time / note
P1	1e65261c	54255188	604f5375	09183973	871.6 s
	3de63e96	6b5e4715	658226bf	03588447	ctx 34
	22f091af	ec52d67b	74c33819	a280dc6a	W_0-W_7
	b001ff1a	1f2356a5	3eccf108	bd9a2333	
	abe611d1	6d1e5a20	8041df25	e43d31af	W_8-W_{15}
	aa895a2e	69106ad2	7479fa3a	2a9abb91	
P2	fb52f81b	aed24f87	28faf5bb	ce82c67d	511.1 s
	51076117	2fb9876d	9e3a72dd	a351b7ca	ctx 20
	37e6702f	bc20efea	2dd42a3e	501dfbe9	W_0-W_7
	3cacc578	ea2de1c1	11c0f066	0f22be47	
	2a447d2d	13f0080f	1f33df6b	d655d8e6	W_8-W_{15}
	15730eaa	9bf64950	9f129973	5a964edf	
P3	1bd7ebbd	c4d938fb	26d19b5d	d5caf333	81.7 s
	de397bd1	c745727b	d5556baf	38ccf977	ctx 3 [†]
	3ce8fba4	e2fb9661	44730c59	e1cf4bc0	W_0-W_7
	e1a18d93	97658983	67efe2a7	ef260ecb	
	d4c6dbe0	13e9388e	95664a59	4d9e248b	W_8-W_{15}
	74137862	664815ac	89eae95a	cd7dbef5	
P4	c55aecaf	d68ac4f5	59bb2b94	0688f3bc	campaign
	1dabe3ff	86a7d034	47e8b11a	637d7506	target 2, ctx 9
	3694ee67	a1d020b9	47f08ea3	dc873e2b	W_0-W_7
	1f41c8b1	39018970	d3a2afb2	23f6e6cf	
	80c46545	f7b2b46c	2e44f333	2effe68b	W_8-W_{15}
	06d80cdf	8525e7d7	7ef008e5	13e56087	
P5	a8241289	b5980304	e9dfd401	b1370200	campaign
	ab32b5e8	cfaff26c	26aabcf	8be89df4	target 20, ctx 7
	3e0a7d67	df0944a3	4bfe156b	bd344ac5	W_0-W_7
	7b7d9587	8984e286	21c995d9	d8e53cdd	
	0a0be79d	e02dda1a	e17e0f4a	e3ff692b	W_8-W_{15}
	8c782631	97dc4811	5ed68656	dfce314f	
P6	09156002	031fed28	99bc67d0	e419b6ff	campaign
	81b9a0dd	4e3fa417	e000817a	db85b143	target 25, ctx 4
	164084a5	f97bd8bd	bff47431	a8e55b29	W_0-W_7
	948d1160	10eb8f7d	d0ea565e	7918b8b9	
	87527a3a	6baa5685	a1fbc5dc	7548adfc	W_8-W_{15}
	be3fb206	b9738756	b4ca33e3	b8354a95	

measured the dilution directly on the *lo* residual, where the exact-value enhancement is known: a $37\times$ enhancement at the exact residual appears as $3.9\times$ at 256-bin resolution, a dilution of $9.5\times$. A hi enhancement of $13\times$ diluted by the same factor would appear as $1.4\times$, shifting the occupancy of one bin from 22.6 to roughly 31—invisible to a χ^2 statistic over 255 degrees of freedom.

Measuring at exact resolution instead shows the hi filter passing at $7\text{--}13\times$ the uniform rate (E1 and E2, §5.6). The correct conclusion is that the coarse-binned test is underpowered, not that the hi residuals are uniform; filter rates in this attack must be estimated at exact residual resolution.

7 Comparison with Prior Work

Table 3: Preimage results on step-reduced SHA-256. “Concrete” means at least one preimage was computed and verified.

Reference	Rounds	Cost	Target	Concrete
Isobe–Shibutani [3]	24	2^{240}	any	no
Aoki et al. [4]	43	$2^{254.9}$	any	no
Khovratovich et al. [5]	45	$2^{255.5}$	any	no
Davydov et al. [11]	17, 18	SAT, seconds	zero digest	yes
Davydov et al. [11]	19	SAT, unsolved ^a	zero digest	partial
Zaikin [12]	19	param. Cube-and-Conquer ^b	not examined ^b	reported
This work	19	$\approx 2^{32}$ per attempt, ≈ 80 attempts	random digests	six

^a Reached only in a weakened form that constrains at most 31 of the 256 output bits. ^b Reported for the compression function itself; full text, cost and target form not examined (see text).

The theoretical attacks reach far more rounds than any practical method, at costs within a few bits of the generic 2^{256} , and have not been executed. The practical line is SAT-based. Davydov et al. [11] use four independent CNF encodings (CBMC, Transalg, and SAT-encoding in Counter-chain, Dot-matrix, Two-operand and Espresso variants) with the Kissat solver at a 5000s limit, find preimages of the 17- and 18-round compression function (18 rounds in 2.38–20.6s depending on encoding), and report 19 rounds as not solved for every encoding tried. Their weakened 19-round problem constrains only a zero prefix of the output: plain Kissat reached a 24-bit prefix, and a 24-hour cube-and-conquer run on 16 cores solved every case except the 32-bit one, i.e. at most 31 of the 256 output bits. The authors conclude that a 19-round preimage is “too difficult a task even if one carefully selects the SAT encoding and uses parallel computing.” Their published benchmark instances fix the all-zero digest as the target.

Zaikin [12] reports that a parameterized Cube-and-Conquer solver inverted the 19-round compression function, the first such result. The weakened problems in that method are the intermediate instances used to tune the solver; the reported result is stated for the compression function itself. We have not been able to examine the article’s full text, so we cannot confirm that all 256 output bits were fixed, nor state its computational cost, and we make no priority claim. What the present work adds is a different method with an explicit cost profile: after a one-time 16 GB precomputation, each attempt is one 2^{32} sweep of about 25s on a single GPU, about one attempt in eighty succeeded in the campaign of §5.7, and six preimages of random digests are published together with the code that found them.

Attack scope. The attack covers 19 of 64 rounds and targets the compression function with the standard chaining value; the padded hash is discussed in §1. Full SHA-256 is not threatened, and nothing here bears on its security margin.

8 Conclusion

Three cancellation identities in the SHA-256 round and schedule equations, together with a provisional assignment for the part of W_9 that depends on a_2, a_3 , reduce the recovery of a 19-round preimage to one 2^{32} sweep of table lookups per attempt. Six preimages of random digests were found and verified, and a preregistered 240-attempt campaign measured the success rate directly, $\hat{\mu} = 1.25 \times 10^{-2}$ per attempt with a wide interval. A cheap screening pass improves the intermediate yield by a measured $3.41\times$; whether that survives to the final success rate is not established.

The same construction does not reach 20 rounds. The fourth schedule constraint adds a 32-bit condition, and the companion note [1] shows why no rearrangement of this attack family absorbs it: the full-avalanche dependency graph among the absorbed state words acquires its first cycle at exactly 20 rounds, through the term $\Sigma_0(a_4)$ entering W_9 , so the 32-bit iteration that closes the 19-round system has no cheap counterpart. We do not claim that this is fundamental.

Artefacts. Solver, table construction, campaign driver, raw campaign data and an independent verifier are available at

<https://github.com/kamb-code/sha256-r19-preimage>. Any of the six preimages (Table 2) can be checked with

```
python3 code/verify_r19.py --rounds 19 --hash <hex> --words "<W0 ...W15>"
```

A An Exact Cancellation in the Message Schedule

The σ_0 analysis above also exposes a two-word message difference that the schedule cancels exactly for eight consecutive positions.

Lemma 1 (Nine-step cancellation). *Let $W_0, \dots, W_{15}$ be any message block and let $\delta \neq 0$. Define $W'_i = W_i$ for $i \notin \{0, 9\}$, $W'_0 = W_0 + \delta$ and $W'_9 = W_9 - \delta$ (all arithmetic mod 2^{32}). Then the extended schedules satisfy $W'_r - W_r = 0$ for $r = 16, \dots, 23$, and $W'_{25} - W_{25} = -\delta$, for every block and every δ .*

Proof. The schedule recurrence is $W_r = \sigma_1(W_{r-2}) + W_{r-7} + \sigma_0(W_{r-15}) + W_{r-16}$. For $r \in [16, 23]$ the term $\sigma_0(W_{r-15})$ indexes $r - 15 \in [1, 8]$, where the difference is zero; the term $\sigma_1(W_{r-2})$ indexes $r - 2 \in [14, 21]$, where the difference is zero up to $r = 23$ by induction; the term W_{r-7} indexes $r - 7 \in [9, 16]$ and contributes $-\delta$ only at $r = 16$; and the term W_{r-16} indexes $r - 16 \in [0, 7]$ and contributes $+\delta$ only at $r = 16$. At $r = 16$ the two contributions cancel, and for $r \in [17, 23]$ all four terms contribute zero. At $r = 25$ the only nonzero term is $W_{r-16} = W_9$, giving $-\delta$. $\square$

An exhaustive check over all 240 ordered pairs of positions in $[0, 15]$ shows that $(0, 9)$ and $(9, 0)$ are the only pairs for which ΔW_r vanishes identically for all $r \in [16, 23]$ (script `alt_differential.py`). The reason is that $\sigma_0(W_k)$ first enters the schedule at W_{k+15} , so a perturbed position $k \in [1, 8]$ seeds a nonzero difference inside the window, and cancellation of the two linear terms at $r = 16$ then requires exactly the pair $\{0, 9\}$.

Why no search result is claimed. GPU sweeps built on this difference located message pairs with one equal output word of the full 64-round compression function, and multi-block coordinate descent reached output differences of Hamming weight $72/256$. We record these only to state why they are not results. For a fixed input difference, the expected minimum Hamming weight of a random 256-bit output difference over 2^{32} trials is 78, and over the $2^{37} - 2^{39}$ evaluations used by the multi-block searches it is $72\text{--}74$. The observed minima ($75\text{--}95$ over single 2^{32} sweeps, 72 after multi-block descent) are therefore at the level expected of an unbiased

difference, and one-word equalities occurred at the birthday rate of 2^{-32} per output word. The code and data remain in the repository for completeness.

References

- [1] K. S. Basra. The 20-round barrier for table-based preimage attacks on SHA-256. Companion note, 2026. Source and code at <https://github.com/kamb-code/sha256-r19-preimage>.
- [2] National Institute of Standards and Technology. *FIPS PUB 180-4: Secure Hash Standard (SHS)*, 2015.
- [3] T. Isobe and K. Shibutani. Preimage attacks on reduced Tiger and SHA-2. In *Fast Software Encryption (FSE 2009)*, LNCS 5665, pp. 139–155. Springer, 2009.
- [4] K. Aoki, J. Guo, K. Matusiewicz, Y. Sasaki, and L. Wang. Preimages for step-reduced SHA-2. In *Advances in Cryptology – ASIACRYPT 2009*, LNCS 5912, pp. 578–597. Springer, 2009.
- [5] D. Khovratovich, C. Rechberger, and A. Savelieva. Bicliques for preimages: attacks on Skein-512 and the SHA-2 family. In *Fast Software Encryption (FSE 2012)*, LNCS 7549, pp. 244–263. Springer, 2012.
- [6] S. K. Sanadhya and P. Sarkar. New collision attacks against up to 24-step SHA-2. In *Progress in Cryptology – INDOCRYPT 2008*, LNCS 5365, pp. 91–103. Springer, 2008.
- [7] F. Mendel, T. Nad, and M. Schl  ffer. Improving local collisions: new attacks on reduced SHA-256. In *Advances in Cryptology – EUROCRYPT 2013*, LNCS 7881, pp. 262–278. Springer, 2013.
- [8] Y. Li, F. Liu, and G. Wang. New records in collision attacks on SHA-2. In *Advances in Cryptology – EUROCRYPT 2024*, LNCS. Springer, 2024. DOI 10.1007/978-3-031-58716-0_6.
- [9] Y. Li, F. Liu, G. Wang, X. Dong, and S. Sun. The first practical collision for 31-step SHA-256. In *Advances in Cryptology – ASIACRYPT 2024*, LNCS. Springer, 2024. DOI 10.1007/978-981-96-0941-3_8.
- [10] I. Mironov and L. Zhang. Applications of SAT solvers to cryptanalysis of hash functions. In *Theory and Applications of Satisfiability Testing (SAT 2006)*, LNCS 4121, pp. 102–115. Springer, 2006.
- [11] V. V. Davydov, M. D. Pikhtovnikov, A. P. Kiryanova, and O. S. Zaikin. Analysis of the cryptographic strength of the SHA-256 hash function using the SAT approach. *Scientific and Technical Journal of Information Technologies, Mechanics and Optics*, 25(3):428–437, 2025. DOI 10.17586/2226-1494-2025-25-3-428-437.
- [12] O. Zaikin. Preimage attacks on round-reduced MD5, SHA-1, and SHA-256 using parameterized SAT solver. *Constraints*, published online 23 January 2026. DOI 10.1007/s10601-025-09383-0.